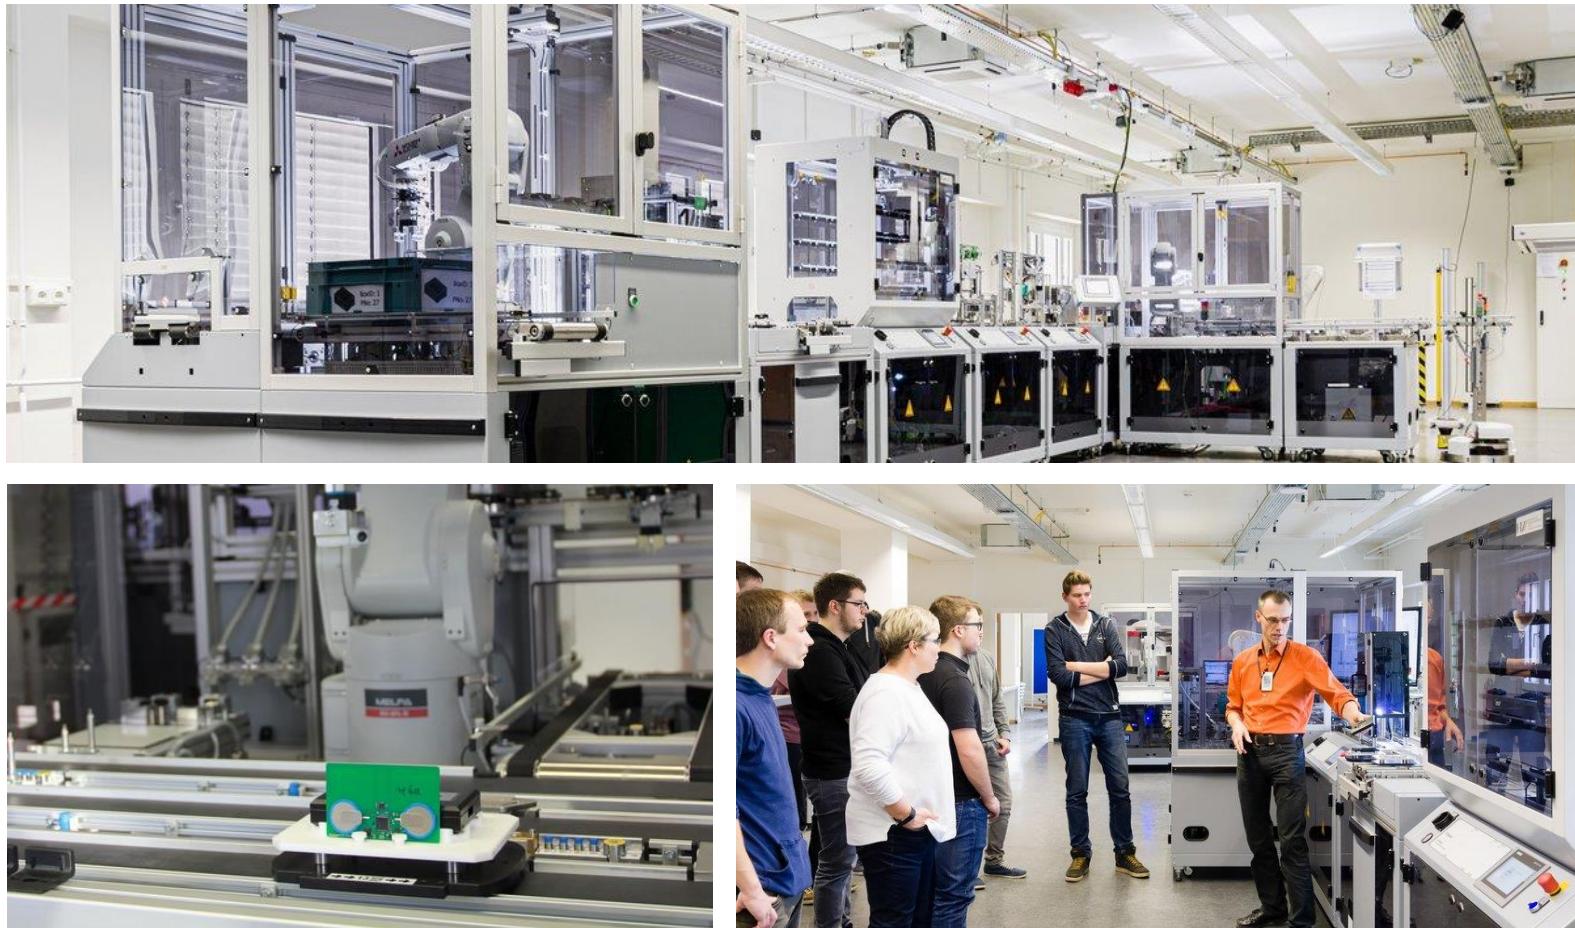


# Anomaly Detection with Deep Learning

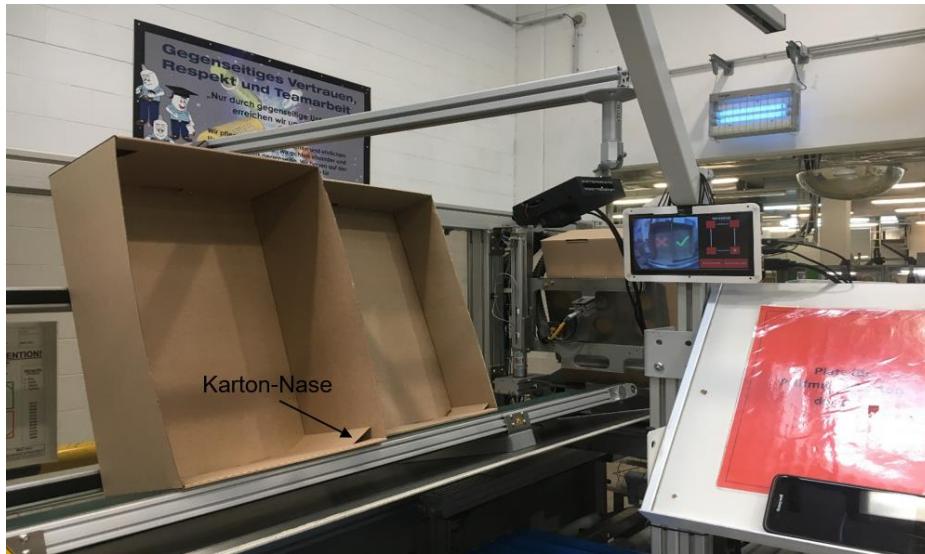


You can do  
deep learning

- Johannes Metzler, research associate in research group **"Smart Production System"** from Prof. Dirk Reichelt
- Expert in **Computer Vision** and **Artificial Intelligence**



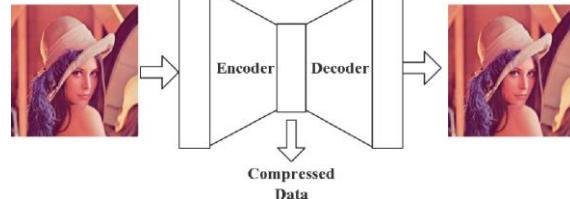
# Projects done.



# What we do in this course?

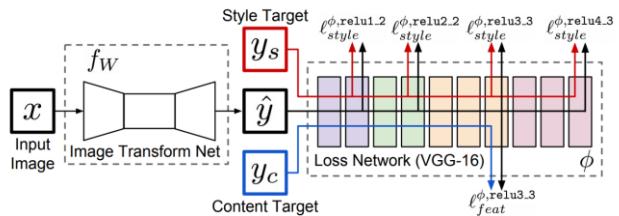


Foundations of Computer Vision



Anomaly Detection with  
Autoencoders

Boosting Autoencoders with  
Feature Loss



State of the art Anomaly  
Detection with Anomalib



Black box

- Interpretable ML
- Visualize gradients and activations

Needs too much data

- Transfer learning
- Share pre-trained nets

Needs ML PhD

- No longer true
- fastai & keras libs, MOOCs, etc

Only for vision

- No longer true
- SoTA for speech, structured data, time series...

Needs lots of GPUs

- Was never true
- ...except for some research projects

“Not really AI”

- Who cares?
- Do you really want to build a brain?

Code first

Focus on learning from experiments

The whole game

It's like learning soccer as a kid (Perkins)

Concepts, not details

We'll gradually dig in to all the details

You will learn how it works before you learn why it works!

# Where is deep learning the best known approach?

NLP	answering questions; speech recognition; summarizing documents; classifying documents; finding names, dates, etc. in documents; searching for articles mentioning a concept
Computer vision	satellite and drone imagery interpretation (e.g. for disaster resilience); face recognition; image captioning; reading traffic signs; locating pedestrians and vehicles in autonomous vehicles
Medicine	finding anomalies in radiology images, including CT, MRI, and x-ray; counting features in pathology slides; measuring features in ultrasounds; diagnosing diabetic retinopathy
Biology	folding proteins; classifying proteins; many genomics tasks, such as tumor-normal sequencing and classifying clinically actionable genetic mutations; cell classification; analyzing protein/protein interactions
Image generation	colorizing images; increasing image resolution; removing noise from images; converting images to art in the style of famous artists
Recommendation systems	web search; product recommendations; home page layout
Playing games	better than humans and better than any other computer algorithm at Chess, Go, most Atari videogames, many real-time strategy games
Robotics	handling objects that are challenging to locate (e.g. transparent, shiny, lack of texture) or hard to pick up
Other applications	financial and logistical forecasting; text to speech; much much more...

ImageNet dataset ~1.6 Mio pictures

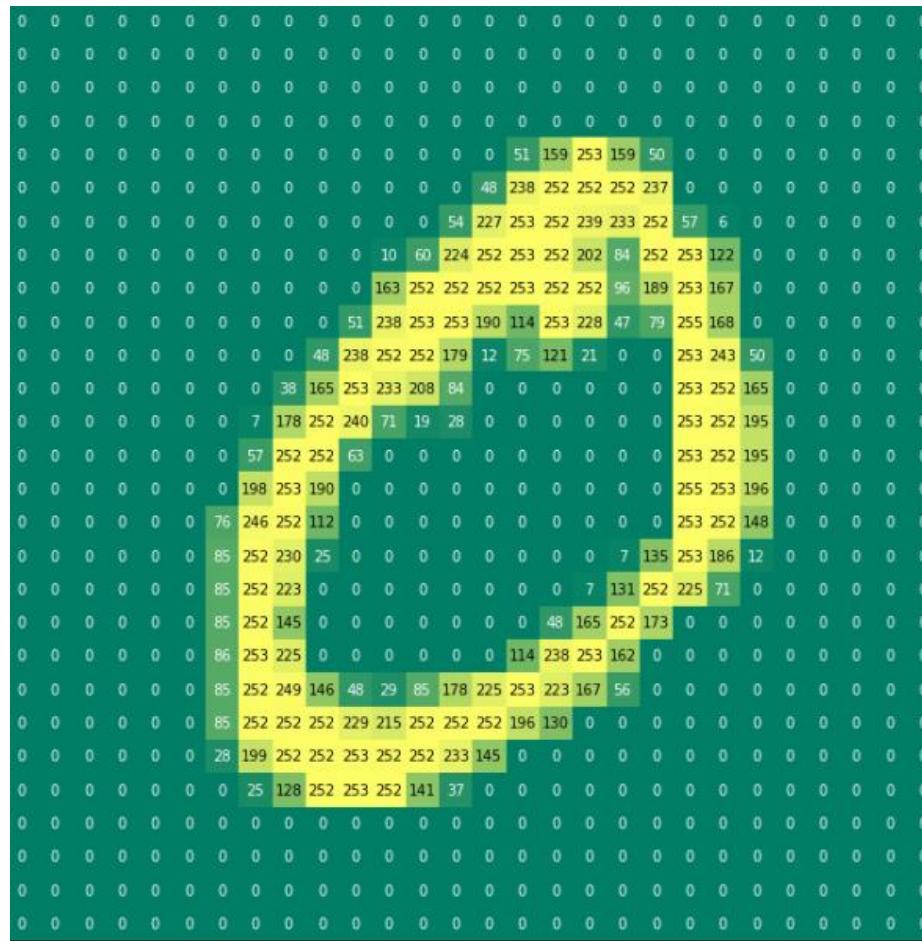
Konfiguration	Float 32 Training	Float 16 Training
CPU(32 cores)	27 Stunden	27 Stunden
Single RTX 2080TI	69 Minuten	29 Minuten
Single RTX 3080	53 Minuten	22 Minuten
Single RTX 3090	41 Minuten	18 Minuten
Single Tesla A100	23 Minuten	8,5 Minuten
4 x RTX 2080TI	19 Minuten	8 Minuten
4 x Tesla V100	15 Minuten	4,5 Minuten
4 x RTX 3090	11,5 Minuten	5 Minuten
4 x Tesla A100	6,5 Minuten	3 Minuten

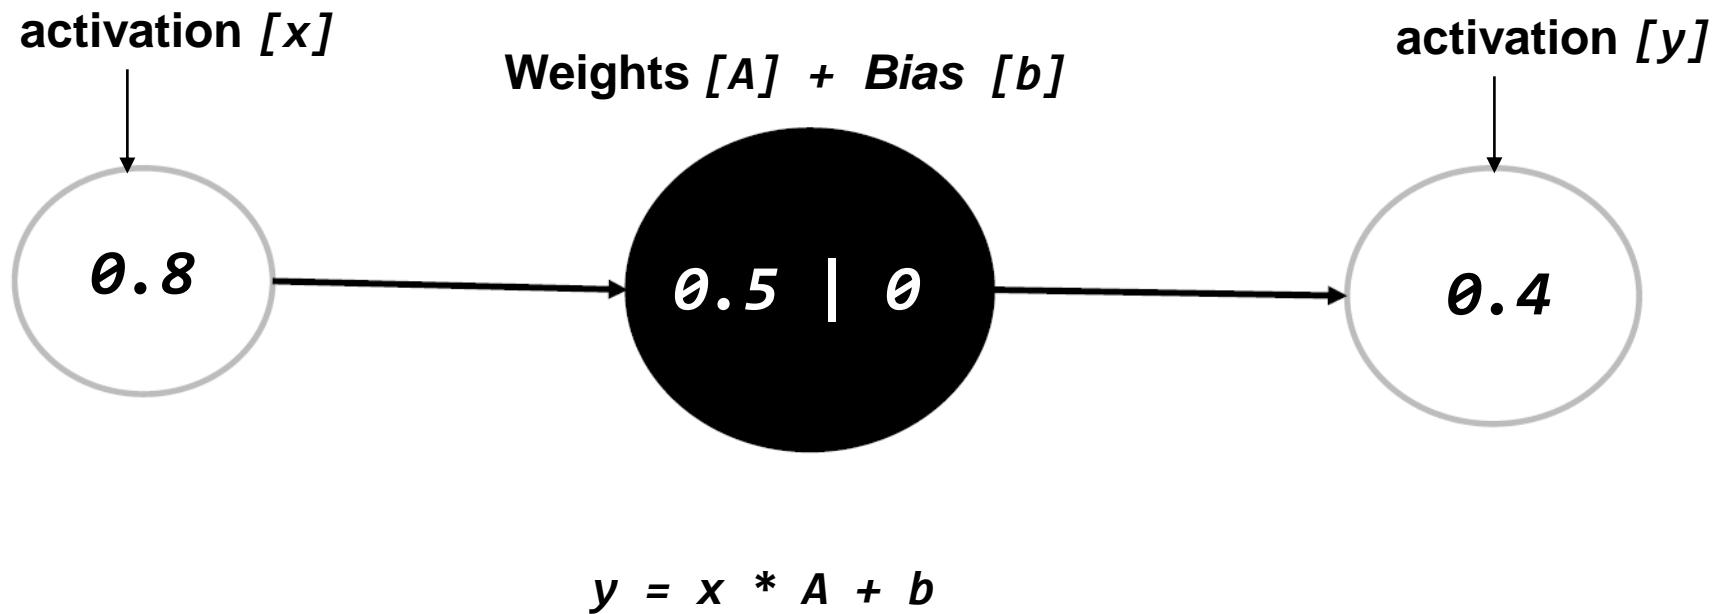
---

# Lab session

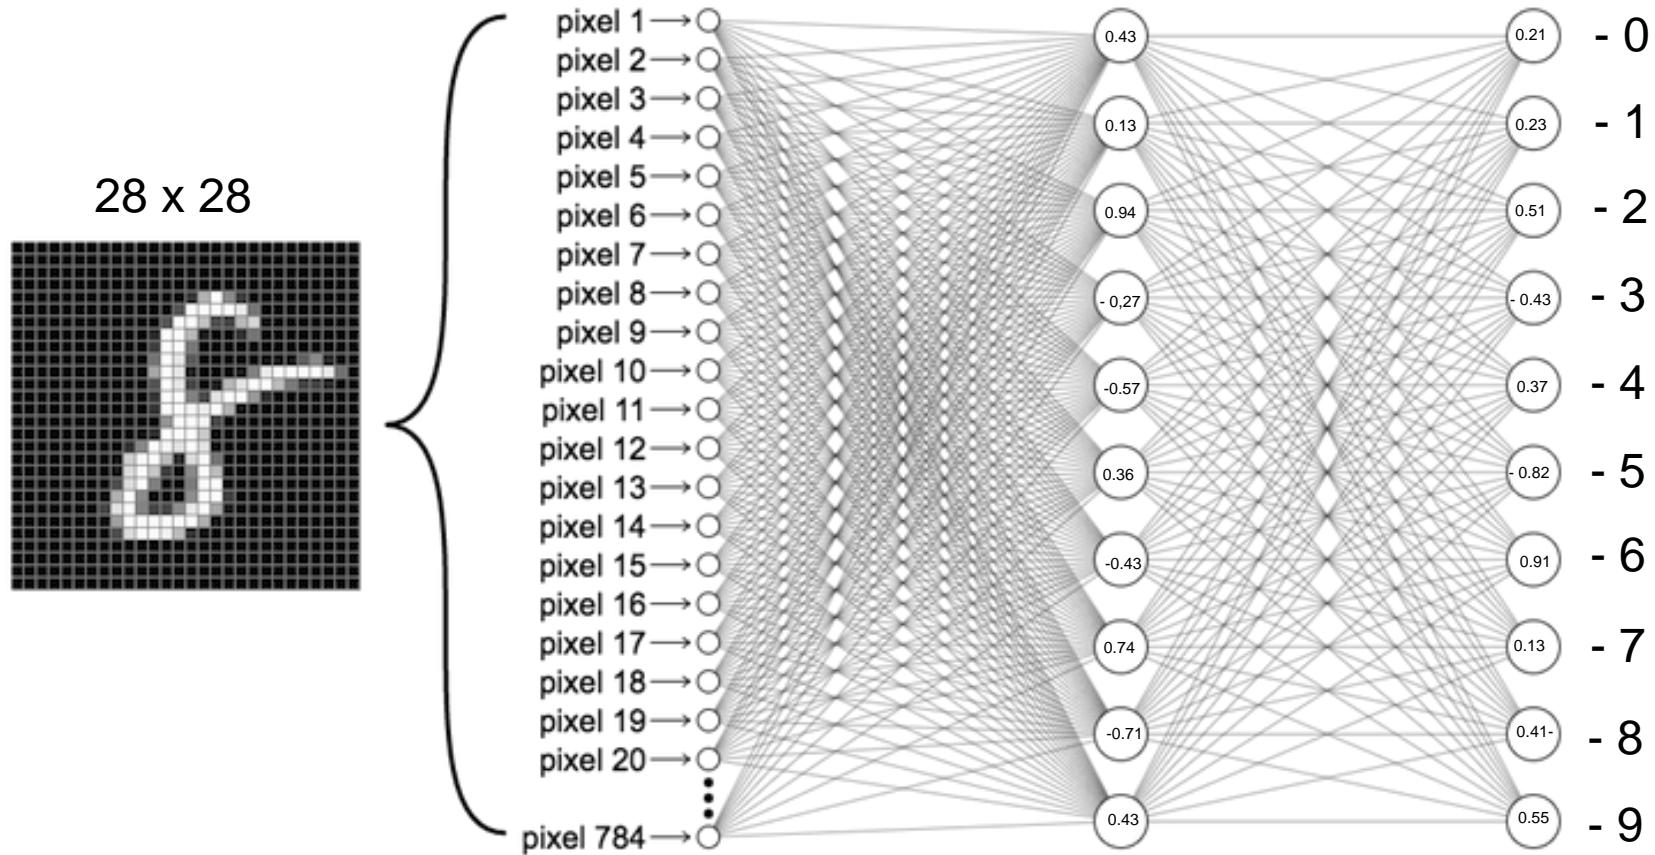


# How a computer sees numbers

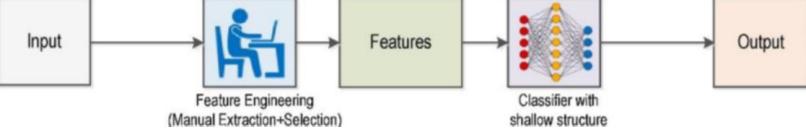
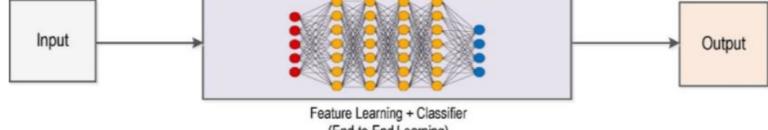


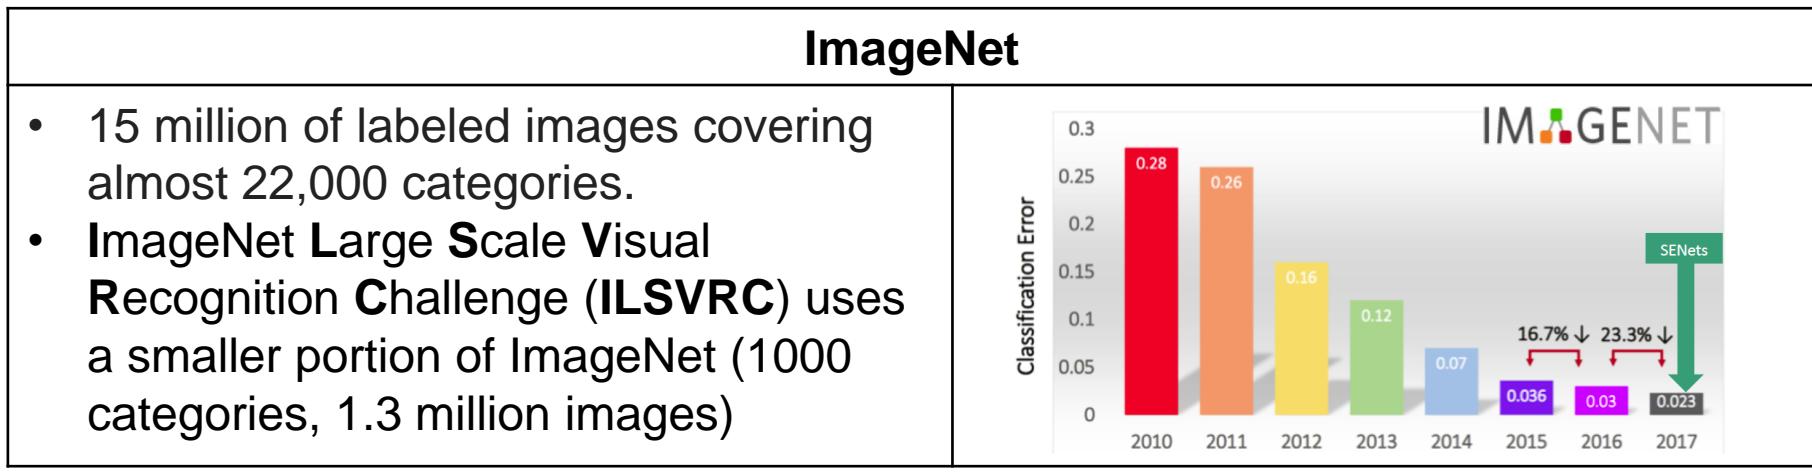


## Forward Propagation



## Backward Propagation

Traditional CV / ML	Deep Learning
	
<ul style="list-style-type: none"> <li>• Works better on small data</li> <li>• Financially and computationally cheap</li> <li>• Easier to interpret</li> </ul>	<ul style="list-style-type: none"> <li>• Best-in-class performance</li> <li>• Scales effectively with data</li> <li>• No need for feature engineering</li> <li>• Adaptable and transferable</li> </ul>



## Supervised Learning

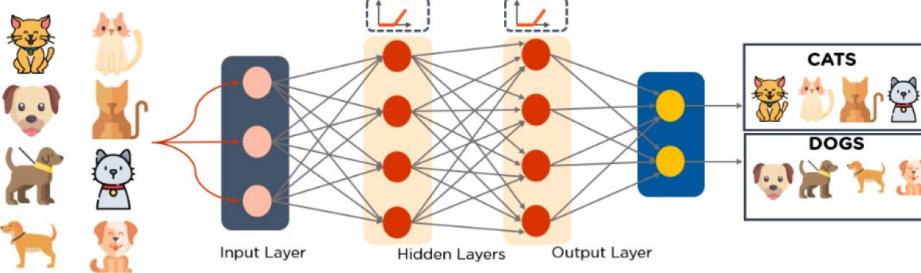
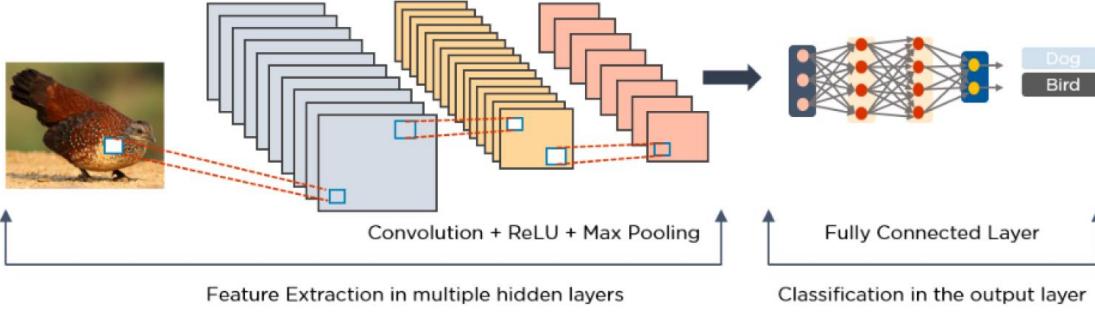
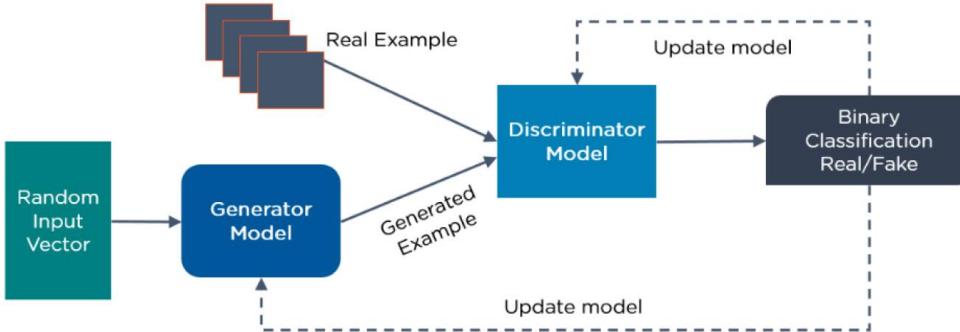
- Regression
  - *Predictions*
  - *Optimizations*
  - *Modeling*
  - *Forecasts*
- Classification
  - *Diagnosis*
  - *Sorting*
  - *Recommender System*
- *Algorithms*
  - *Linear Regression*
  - *Logistic Regression*
  - *Random Forests*
  - *Multilayer Percept.*
  - *CNN*
  - *RNN*
  - .....

## Unsupervised Learning

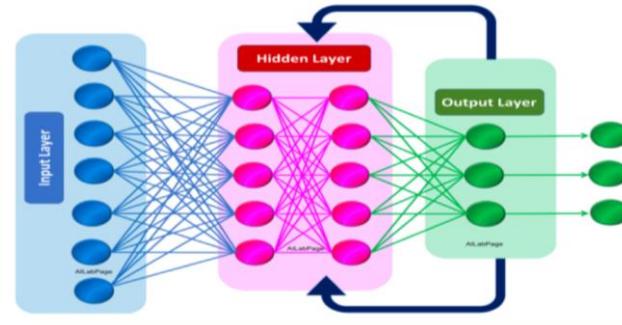
- Clustering
  - *Grouping*
  - *Pattern recognition*
- Dimension Reduction
- Anomaly Detection
- Algorithms
  - K-Means
  - PCA
  - Gaussian Mixture
  - Singular value decomposition
  - Autoencoder

## Reinforcement Learning

- Game Intelligence (Alpha Go)
- Swarm Intelligence
- Machine optimization
- Algorithms
  - Monte Carlo
  - Deep Q-learning
  - State–action–reward–state–action (SARSA)

1. Multilayer Perceptrons (MLPs)	
2. Convolutional Neural Networks (CNNs)	 <p>Convolution + ReLU + Max Pooling</p> <p>Feature Extraction in multiple hidden layers</p> <p>Fully Connected Layer</p> <p>Classification in the output layer</p>
3. Generative Adversarial Networks (GANs)	 <p>Random Input Vector</p> <p>Generator Model</p> <p>Discriminator Model</p> <p>Binary Classification Real/Fake</p> <p>Update model</p> <p>Generated Example</p> <p>Real Example</p>

## 4. Recurrent Neural Networks (RNN)



## 5. Transformer Networks (TN)



# Convolutional Neural Networks (CNN) 1/4

2012 to Present: CNN's are everywhere

Image Classification



Object Detection

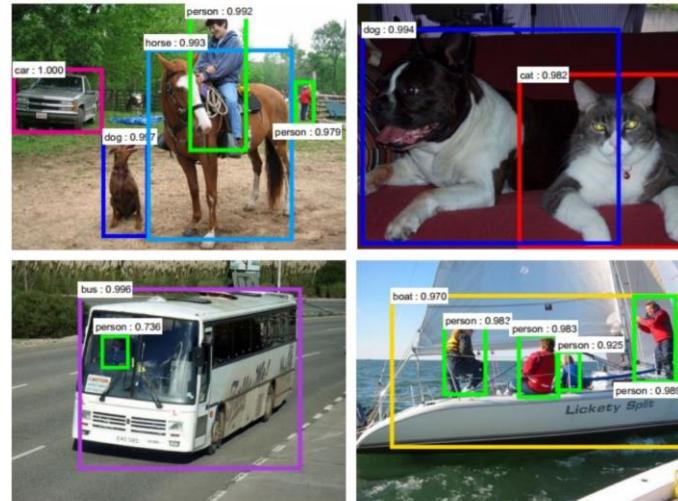


Image Retrieval



Image Segmentation



2012 to Present: CNN's are everywhere

## Pose Recognition



## Image Captioning



*A white teddy bear  
sitting in the grass*



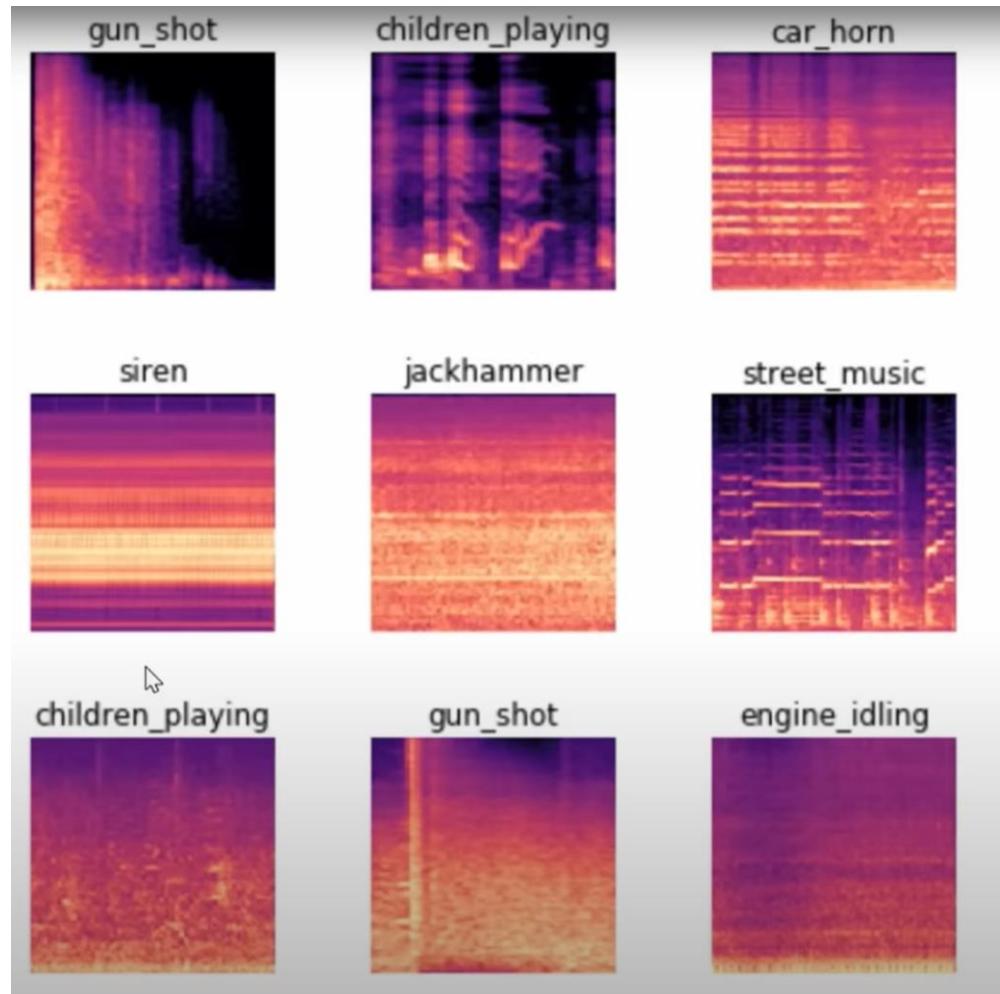
*A man in a baseball  
uniform throwing a ball*

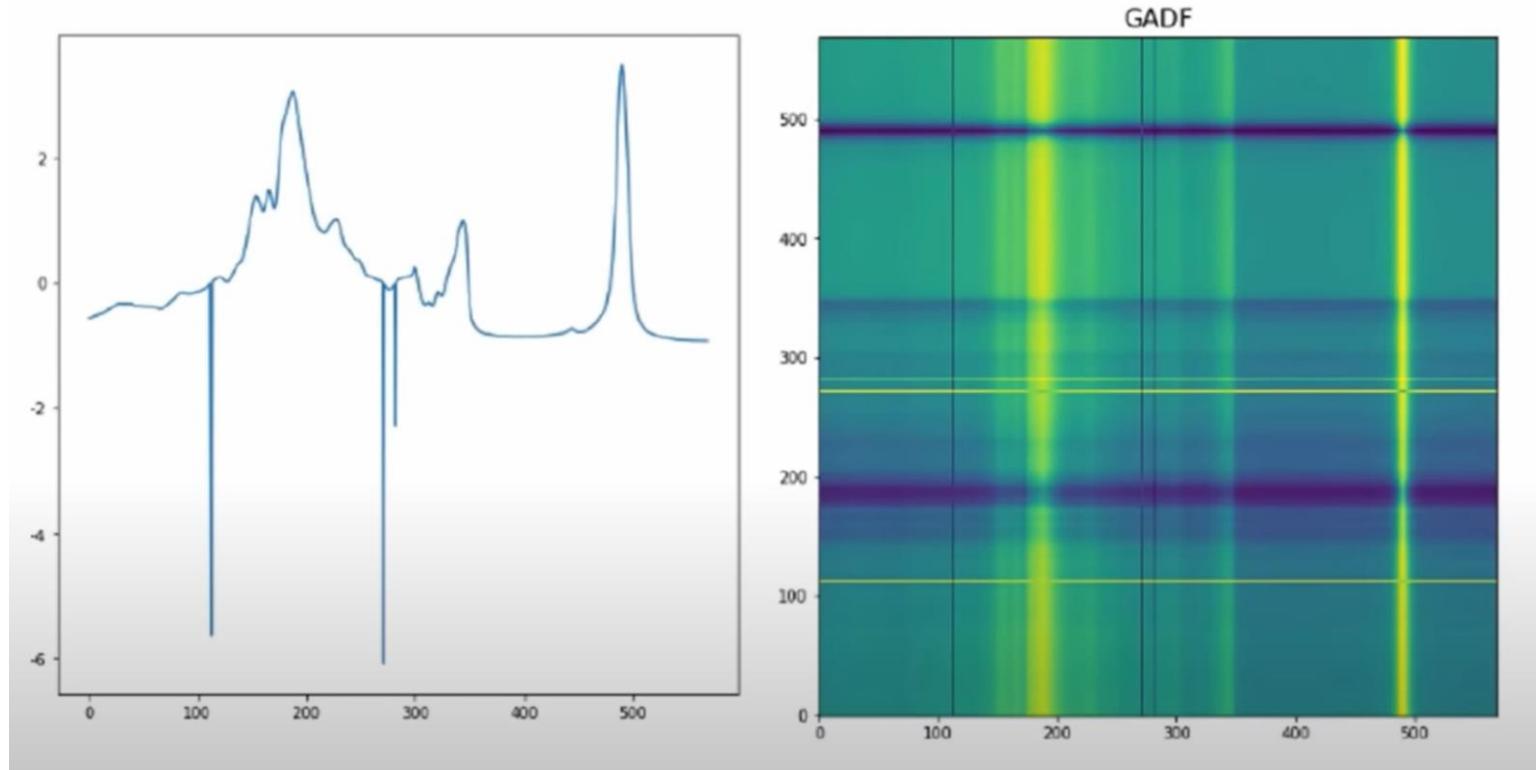


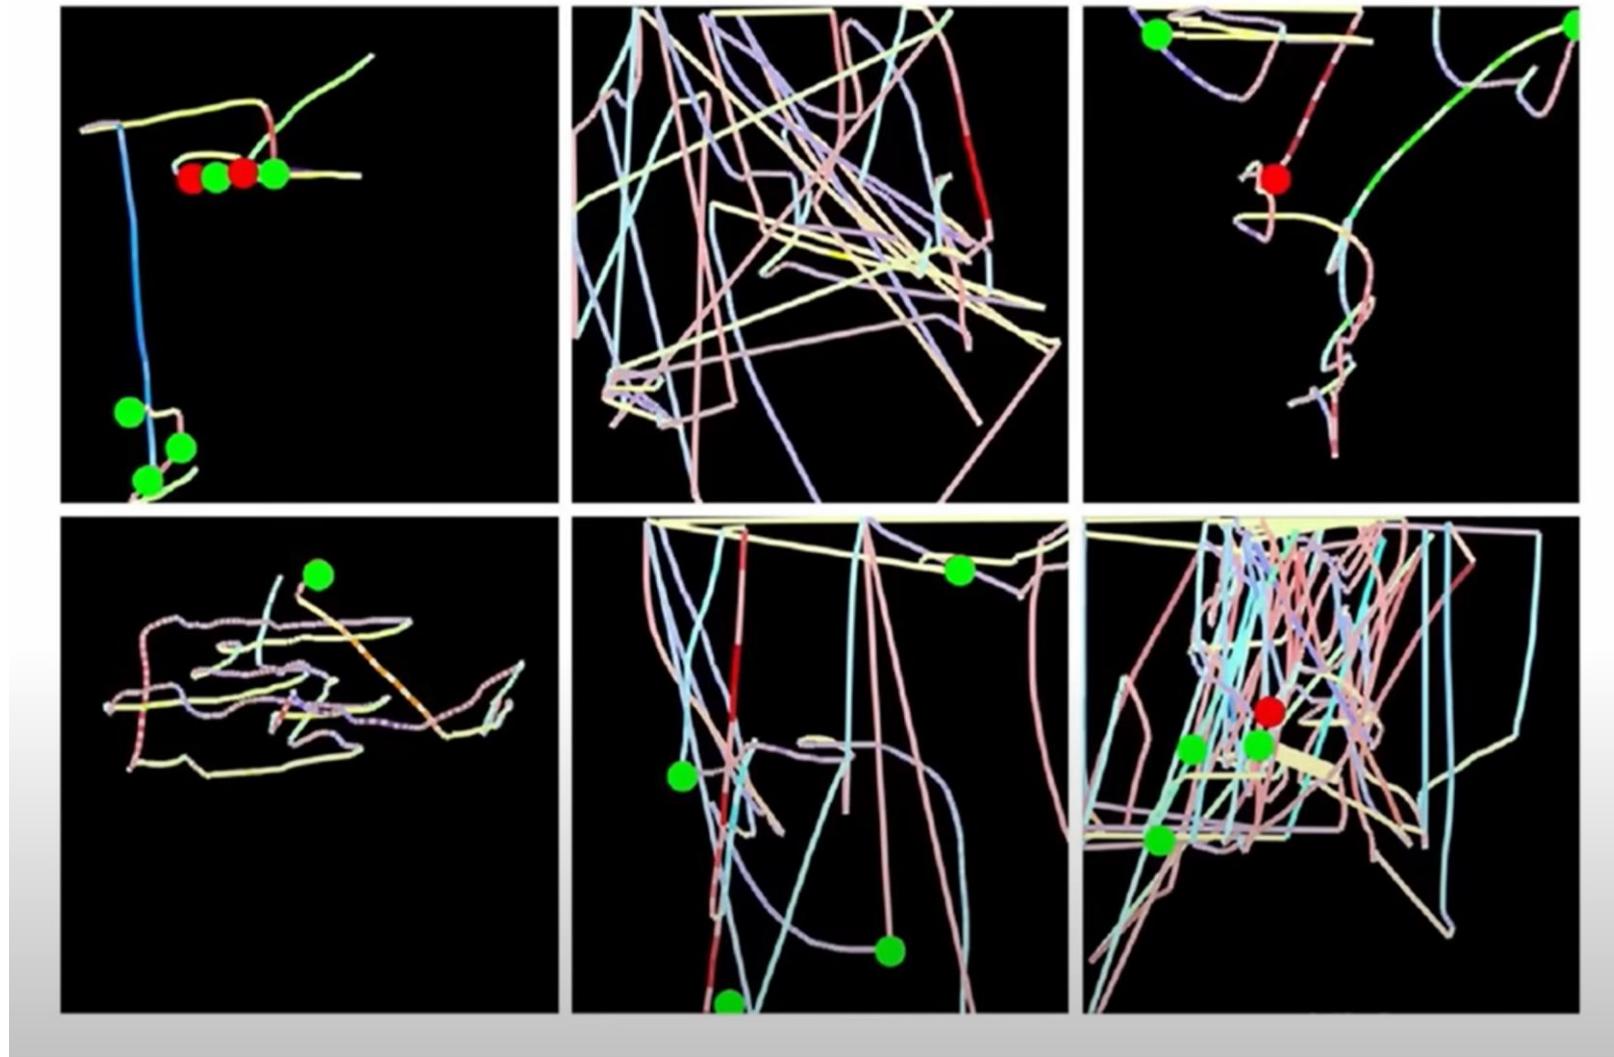
*A woman is holding  
a cat in her hand*

2012 to Present: CNN's are everywhere

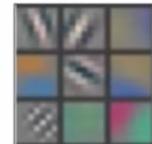




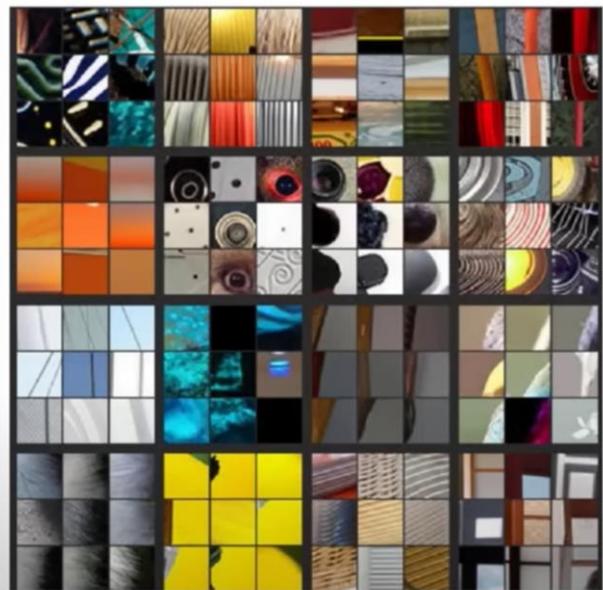
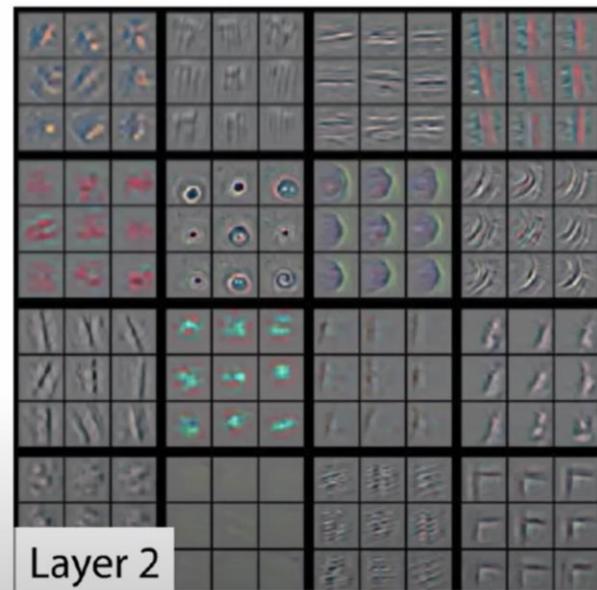




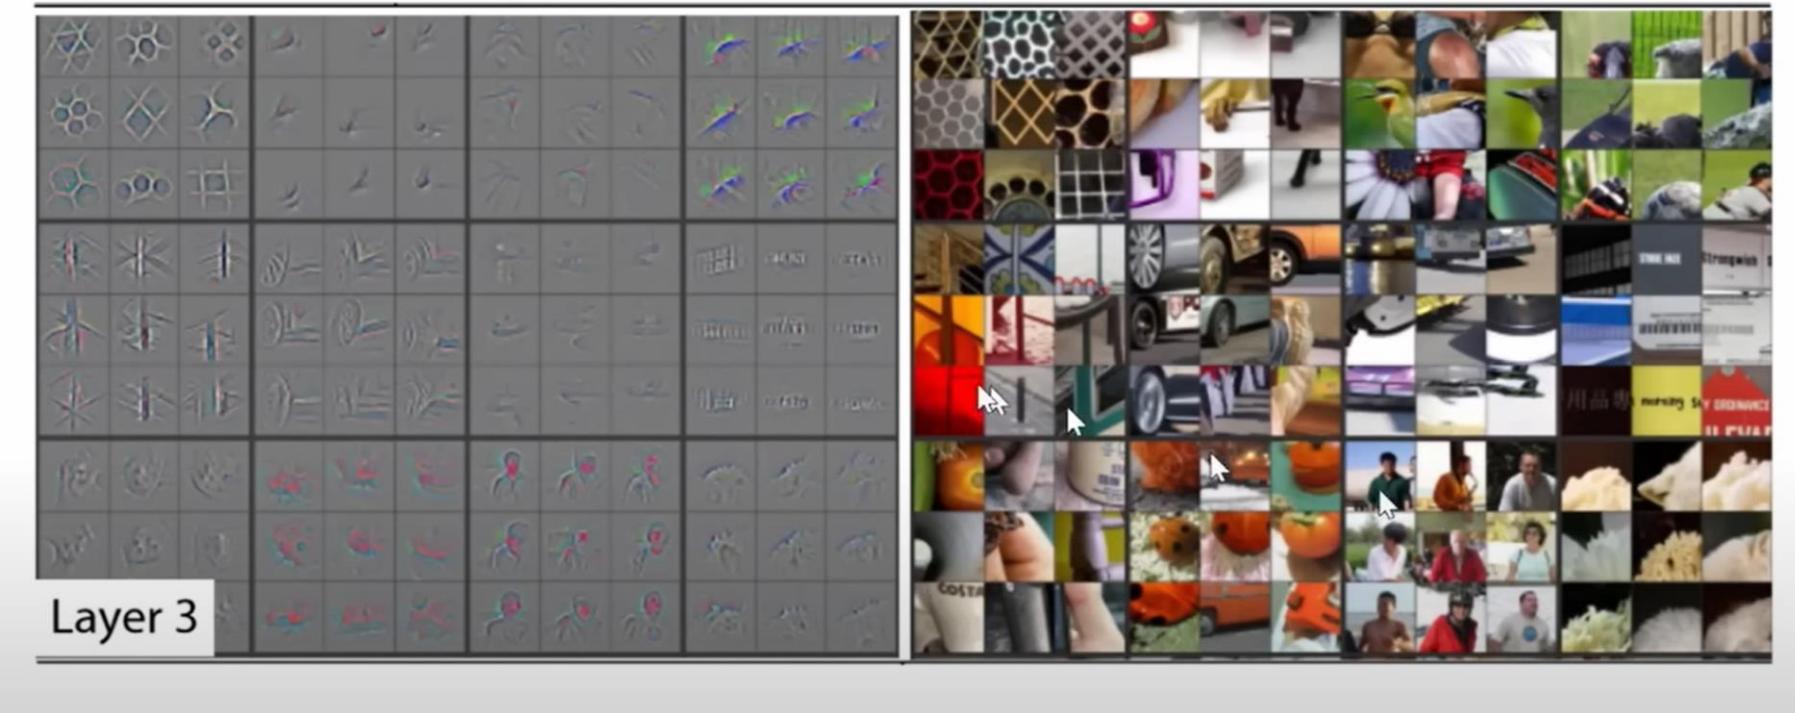
# How CNN works



Layer 1



# How CNN works



Lesson 1: Practical Deep Learning for Coders 2022

Industry Dec 14, 2021

## PyTorch vs TensorFlow in 2022

Should you use PyTorch vs TensorFlow in 2022? This guide walks through the major pros and cons of PyTorch vs TensorFlow, and how you can pick the right framework.

Ryan O'Connor  
Developer Educator

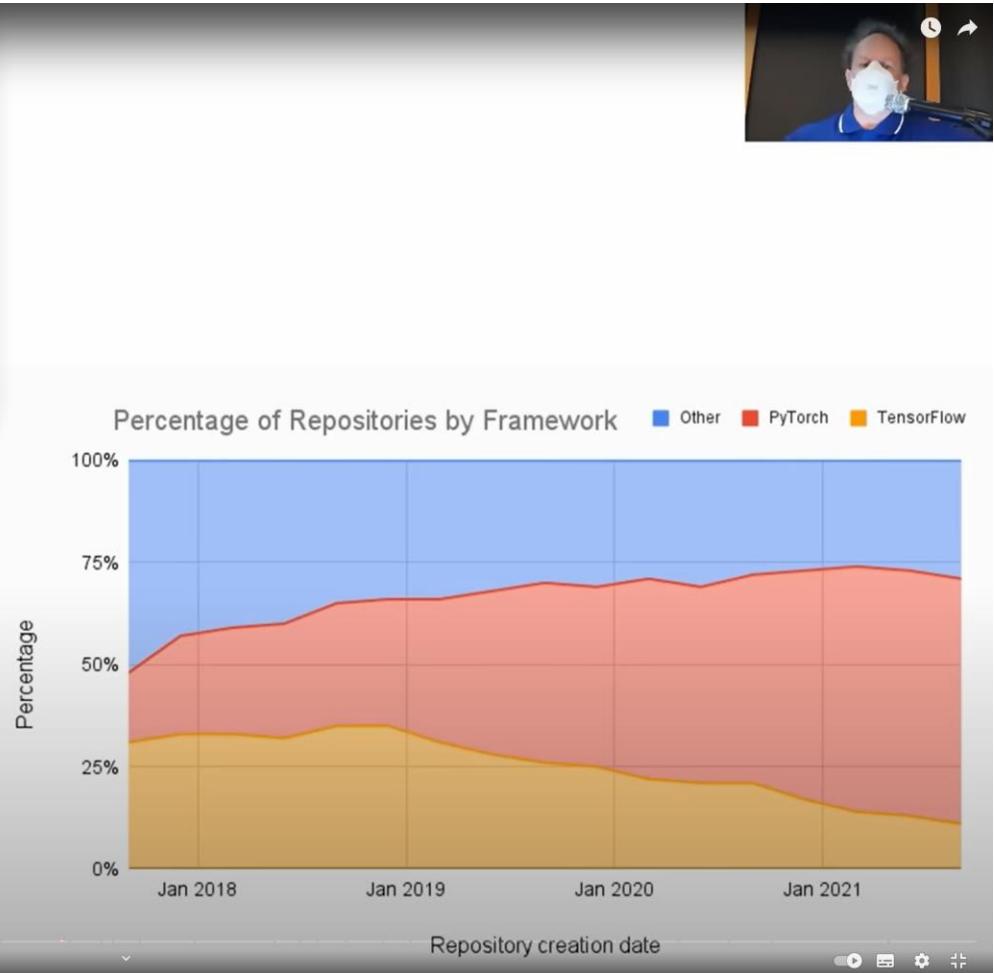


The left diagram shows the percentage of repositories for TensorFlow and PyTorch in 2018 and 2019. In 2018, PyTorch was at 15% and TensorFlow at 85%. By 2019, TensorFlow had dropped to 45% while PyTorch had risen to 55%.

Framework	2018 (%)	2019 (%)
TensorFlow	85%	45%
PyTorch	15%	55%

The right diagram is a stacked area chart titled "Percentage of Repositories by Framework". The Y-axis represents the percentage from 0% to 100%. The X-axis represents the repository creation date from Jan 2018 to Jan 2021. The chart shows three categories: TensorFlow (yellow), PyTorch (red), and Other (blue). TensorFlow starts at approximately 30% in Jan 2018 and decreases to about 10% by Jan 2021. PyTorch starts at approximately 20% in Jan 2018 and increases to about 70% by Jan 2021. Other frameworks account for the remaining 50%.

Repository creation date	TensorFlow (%)	PyTorch (%)	Other (%)
Jan 2018	30	20	50
Jan 2019	30	40	30
Jan 2020	20	50	30
Jan 2021	10	70	20



Here's AdamW's  
step in PyTorch.

```
loss = None
if closure is not None:
    loss = closure()

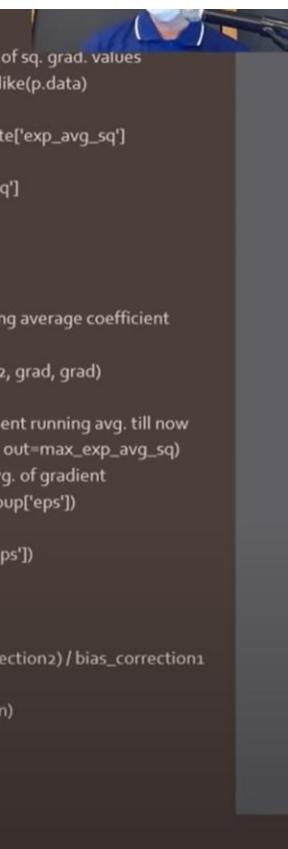
for group in self.param_groups:
    for p in group['params']:
        if p.grad is None:
            continue

        # Perform stepweight decay
        p.data.mul_(1 - group['lr'] * group['weight_decay'])

        # Perform optimization step
        grad = p.grad.data
        if grad.is_sparse:
            raise RuntimeError('Adam does not support sparse gradients, please'
                               'consider SparseAdam instead')
        amsgrad = group['amsgrad']

        state = self.state[p]

        # State initialization
        if len(state) == 0:
            state['step'] = 0
            # Exponential moving average of gradient values
            state['exp_avg'] = torch.zeros_like(p.data)
            # Exponential moving average of squared gradient values
            state['exp_avg_sq'] = torch.zeros_like(p.data)
        if amsgrad:
```



```
# Maintains max of all exp. moving avg. of sq. grad. values
state['max_exp_avg_sq'] = torch.zeros_like(p.data)

exp_avg, exp_avg_sq = state['exp_avg'], state['exp_avg_sq']
if amsgrad:
    max_exp_avg_sq = state['max_exp_avg_sq']
    beta1, beta2 = group['betas']

    state['step'] += 1

    # Decay the first and second moment running average coefficient
    exp_avg.mul_(beta1).add_(1 - beta1, grad)
    exp_avg_sq.mul_(beta2).addcmul_(1 - beta2, grad, grad)
if amsgrad:
    # Maintains the maximum of 2nd moment running avg. till now
    torch.max(max_exp_avg_sq, exp_avg_sq, out=max_exp_avg_sq)
    # Use the max. for normalizing running avg. of gradient
    denom = max_exp_avg_sq.sqrt().add_(group['eps'])
else:
    denom = exp_avg_sq.sqrt().add_(group['eps'])

bias_correction1 = 1 - beta1 ** state['step']
bias_correction2 = 1 - beta2 ** state['step']
step_size = group['lr'] * math.sqrt(bias_correction2) / bias_correction1

p.data.addcdiv_(-step_size, exp_avg, denom)

return loss
```

# Why we use Pytorch and FastAI



```
def weight_decay(p, lr, wd, do_wd=True, **kwargs):
    "Weight decay as decaying `p` with `lr*wd`"
    if do_wd: p.data.mul_(1 - lr*wd)
    return p
weight_decay.defaults = dict(wd=0.)
```

```
"A `Optimizer` for Adam with `lr`, `mom`, `sqr_mom`, `eps` and `params`"
steppers = [] if wd==0. else [weight_decay] if decouple_wd else [l2_reg]
steppers.append(adam_step)
stats = [partial(average_grad, dampening=True), average_sqr_grad, step_stat]
return Optimizer(params, steppers, stats=stats, lr=lr, mom=mom, sqr_mom=sqr_mom, eps=eps, wd=wd)
```

In fastai it's very different

# Lab session